

Курсовая работа на тему "Исследование датасетов"

Выполнена студентом группы М80-114М-22 Дмитроченко Б.А.

0. Выбор датасета

Для курсового проекта был выбран датасет Stanford Cars (<https://paperswithcode.com/datasetstanford-cars>)

1. Импорт датасета

```
In [1]: !pip install opendatasets --upgrade --quiet

In [2]: import os
import torch
import torchvision
import tarfile
import torch.nn as nn
import numpy as np
import cv2
import torchvision.transforms.functional as F
from torchvision.datasets import ImageFolder
from torch.utils.data import DataLoader
import torchvision.transforms as T
from torch.utils.data import random_split
from torchvision.utils import make_grid
from torchvision.transforms import ToTensor
import matplotlib
import matplotlib.pyplot as plt
matplotlib.interactive(True)

matplotlib.rcParams['figure.facecolor']='w'

C:\Users\boris\jupyter\envs\conda3\lib\site-packages\tqdm\auto.py:22: TqdmWarning: IPProgress not found. Please update jupyter and ipynbwidgets. See https://ipywidgets.readthedocs.io/en/stable/user_install.html
from .autonotebook import tqdm as notebook_tqdm

In [3]: import opendatasets as od
dataset_url = 'https://www.kaggle.com/juotrera/stanford-car-dataset-by-classes-folder'
od.download(dataset_url)

Skipping, found downloaded files in ".\stanford-car-dataset-by-classes-folder" (use force=True to force download)

In [4]: from torch.utils.data import Dataset

In [5]: import os
DATA_DIR_TRAIN = './stanford-car-dataset-by-classes-folder/car_data/car_data/train'
train_classes = os.listdir(DATA_DIR_TRAIN)
DATA_DIR_TEST = './stanford-car-dataset-by-classes-folder/car_data/car_data/test'
test_classes = os.listdir(DATA_DIR_TEST)
train_classes[5:10], test_classes[5:10]

Out[5]: (['Acura ZDX Hatchback 2012',
'AM General Hummer SUV 2000',
'Aston Martin V8 Vantage Convertible 2012',
'Aston Martin V8 Vantage Coupe 2012',
'Aston Martin Virage Convertible 2012'],
['Acura ZDX Hatchback 2012',
'AM General Hummer SUV 2000',
'Aston Martin V8 Vantage Convertible 2012',
'Aston Martin V8 Vantage Coupe 2012',
'Aston Martin Virage Convertible 2012'])

In [6]: def find_classes(dir):
train_classes = os.listdir(dir)
train_classes.sort()
train_class_to_idx = {}
for train_class in train_classes:
train_class_to_idx[train_class] = 0
for test_class in test_classes:
test_class_to_idx[test_class] = 1
return train_classes, test_class_to_idx

train_classes, train_c_to_idx = find_classes(DATA_DIR_TRAIN)
test_classes, test_c_to_idx = find_classes(DATA_DIR_TEST)

In [7]: len(train_classes), len(test_classes)

Out[7]: (196, 196)

In [8]: def find_classes(dir):
test_classes = os.listdir(dir)
test_classes.sort()
test_class_to_idx = {}
for test_class in test_classes:
test_class_to_idx[test_class] = 1
return test_classes, test_class_to_idx

train_classes, test_c_to_idx = find_classes(DATA_DIR_TEST)

In [9]: def extract_class(datasets, slice):
count = 0
for vals in os.listdir(datasets):
print(vals)
count+=1
if count>=slice:
break;

In [10]: extract_class(DATA_DIR_TRAIN, 10)
Acura Integra Type R 2001
Acura RL Sedan 2012
Acura TL Sedan 2012
Acura TL Type-S 2008
Acura TSX Sedan 2012
Acura ZDX Hatchback 2012
AM General Hummer SUV 2000
Aston Martin V8 Vantage Convertible 2012
Aston Martin V8 Vantage Coupe 2012
Aston Martin Virage Convertible 2012
Aston Martin Virage Coupe 2012

In [11]: train_dataset = ImageFolder(DATA_DIR_TRAIN, transform = ToTensor())
test_dataset = ImageFolder(DATA_DIR_TEST, transform = ToTensor())

In [12]: for _ in range(5):
image, label = train_dataset[np.random.randint(len(train_dataset))]
print('image_shape:', image.shape, 'label:', label)
image_shape: torch.Size([3, 224, 441]) label: 29
image_shape: torch.Size([3, 480, 640]) label: 380
image_shape: torch.Size([3, 225, 380]) label: 94
image_shape: torch.Size([3, 144, 259]) label: 77
image_shape: torch.Size([3, 540, 900]) label: 133

In [13]: for _ in range(5):
image, label = test_dataset[np.random.randint(len(test_dataset))]
print('image_shape:', image.shape, 'label:', label)
image_shape: torch.Size([3, 480, 640]) label: 123
image_shape: torch.Size([3, 749, 1024]) label: 0
image_shape: torch.Size([3, 768, 1024]) label: 188
image_shape: torch.Size([3, 124, 180]) label: 66
image_shape: torch.Size([3, 435, 653]) label: 46
```

2. Описание датасета

Датасет состоит из 196 классов автомобилей с общим количеством изображений в 16185 штур, снятых сзади. Количество тренировочных и тестовых данных примерно равно (8144 тренировочных и 8041 тестовых). Присутствуют категории "Make" (марка производителя), "Model" (модель), "Year" (год выпуска - прим. автора). Размер изображений - 360x240. Впервые датасет был представлен в 2017-м году в статье, посвященной разработке сверточной нейронной сети для распознавания изображений (статья - <https://arxiv.org/pdf/1702.01721.pdf>). Сегодня этот датасет используется для протектирования сверточных нейронных сетей, а также для исследования реальных и сгенерированных изображений в факторизации, предпринимательского пространства объектов, таким образом улучшается качество изображений, повышается эффективность выборки и скорость конвергенции.

Генеративно-состязательные сети нацелены на моделирование распределения данного обучающего набора данных. Генератор G сопоставляет скрытые векторы z , отобранные из простого распределения P_z (обычно из нормального распределения), с соответствующими сгенерированными выборками $G(z)$. Затем дискриминатор D нацелен на то, чтобы отличать реальные выборки $x \sim P_x$ от сгенерированных выборок $G(z) \sim P_G(z)$. Эта основная идея приводит к следующей формуле:

$$\min_G \max_z [E_x[\log D(x)] + E_z[\log(1 - D(G(z)))]]$$

Мы выводим набор протекторов объектов (P_x), которые отображают реальные и сгенерированные изображения во входное пространство дискриминатора. Таким образом, прогнозируемое обучение дискриминаторно-состязательной сети можно сформировать следующим образом

$$\min_G \max_{z \sim P_z} [E_x[\log D(P(x))] + E_z[\log(1 - D(G(z)))]]$$

3. Пример данных с разметкой и загрузка тестового набора

```
In [14]: def show_example(img, label):
print('Label: ', train_dataset.classes[label], "(" + str(label) + ")")
plt.imshow(img.permute(1, 2, 0))

In [15]: show_example(train_dataset[np.random.randint(len(train_dataset))])
Label: Dodge Durango SUV 2012 (91)

In [16]: show_example(test_dataset[np.random.randint(len(test_dataset))])
Label: Nissan 240SX Coupe 1998 (107)

In [17]: train_tfms = T.Compose([T.Resize((256, 256)),
T.RandomRotation(0),
T.ToTensor(),
T.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])
test_tfms = T.Compose([T.Resize((256, 256)),
T.ToTensor(),
T.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])

In [18]: train_dataset = ImageFolder(DATA_DIR_TRAIN, transform = train_tfms)
test_dataset = ImageFolder(DATA_DIR_TEST, transform = test_tfms)

In [19]: test_dataset.classes[42:47]

Out[19]: ['Bentley Continental Supersports Conv. Convertible 2012',
'Bentley Mulsanne Sedan 2011',
'Bugatti Veyron 16.4 Convertible 2009',
'Bugatti Veyron 16.4 Coupe 2009',
'Buick Enclave SUV 2012']

In [20]: len(train_dataset), len(test_dataset)

Out[20]: (8144, 8041)

In [21]: random_seed = 42
torch.manual_seed(random_seed);

In [22]: val_percent = 0.1
val_size = int(val_percent * len(train_dataset))
train_size = len(train_dataset) - val_size
train_ds, val_ds = random_split(train_dataset, [train_size, val_size])
len(train_ds), len(val_ds)

Out[22]: (7339, 814)

In [23]: !pip install jvian --upgrade --quiet

In [24]: jvian.log_dataset(dataset_url = dataset_url, val_size = val_size, random_seed = random_seed)
[jvian] Dataset logged.

In [25]: batch_size = 128

In [26]: train_dl = DataLoader(train_ds, batch_size, shuffle = True, num_workers = 4, pin_memory = True)
val_dl = DataLoader(val_ds, batch_size = 2, num_workers = 4, pin_memory=True)

In [27]: def show_batch(dl):
for images, labels in dl:
fig, ax = plt.subplots(figsize = (12, 6))
ax.set_xticks([]), ax.set_yticks([])
ax.imshow(make_grid(images, nrow=10).permute(1,2,0))
break

In [28]: show_batch(train_dl);

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

In [29]: show_batch(val_dl);

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).
```

4. Пример применения готовой модели на данных

Произведем обучение на сверточной нейронной сети

```
In [30]: simple_model = nn.Sequential(
nn.Conv2d(3, 8, kernel_size=3, stride=1, padding=1),
nn.MaxPool2d(2,2)
)

In [31]: for images, labels in train_dl:
print('images.shape:', images.shape)
out = simple_model(images)
print('out.shape:', out.shape)
break

images_shape: torch.Size([128, 3, 256, 256])
out_shape: torch.Size([128, 8, 128, 128])

In [32]: class ImageClassificationBase(nn.Module):
def training_step(self, batch):
images, labels = batch
out = self(images) # Generate predictions
loss = F.cross_entropy(out, labels) # Calculate loss
return loss

def validation_step(self, batch):
images, labels = batch
out = self(images) # Generate predictions
loss = F.cross_entropy(out, labels) # Calculate loss
acc = accuracy(out, labels) # Calculate accuracy
return ('val_loss': loss.detach(), 'val_acc': acc)

def validation_epoch_end(self, outputs):
batch_losses = [x['val_loss'] for x in outputs] # Combine losses
epoch_acc = [x['val_acc'] for x in outputs] # Combine accuracies
return ('val_loss': epoch_loss.mean(), 'val_acc': epoch_acc.mean())

def epoch_end(self, epoch, result):
print("Epoch [%d], train_loss: {:.4f}, val_loss: {:.4f}, val_acc: {:.4f}".format(
epoch, result['train_loss'], result['val_loss'], result['val_acc']))

def accuracy(outputs, labels):
preds = torch.max(outputs, dim=1)
return torch.tensor(torch.sum(preds == labels).item() / len(preds))

In [33]: class StanfordCarsModel(ImageClassificationBase):
def __init__(self):
super().__init__()
self.network = nn.Sequential(
nn.Conv2d(3, 8, kernel_size=3, padding=1),
nn.ReLU(),
nn.Conv2d(8, 16, kernel_size=3, stride=1, padding=1),
nn.ReLU(),
nn.MaxPool2d(2, 2),
nn.Conv2d(16, 32, kernel_size=3, stride=1, padding=1),
nn.ReLU(),
nn.Conv2d(32, 32, kernel_size=3, stride=1, padding=1),
nn.ReLU(),
nn.Conv2d(32, 64, kernel_size=3, stride=1, padding=1),
nn.ReLU(),
nn.Conv2d(64, 64, kernel_size=3, stride=1, padding=1),
nn.ReLU(),
nn.MaxPool2d(2, 2),
nn.Flatten(),
nn.Linear(64*32*32, 1024),
nn.ReLU(),
nn.Linear(1024, 512),
nn.ReLU(),
nn.Linear(512, 196))

def forward(self, xb):
return self.network(xb)

In [34]: model = StanfordCarsModel()
model

In [35]: for images, labels in train_dl:
print('images.shape:', images.shape)
out = model(images)
print('out.shape:', out.shape)
print('out[0]:', out[0])
break

images_shape: torch.Size([128, 3, 256, 256])
out_shape: torch.Size([128, 196])
out[0]: tensor([ 2.1798e-02, 1.5284e-02, 4.6259e-03, -3.9544e-02, -3.8873e-03,
-2.8839e-02, 4.8956e-02, -1.4946e-02, 8.3471e-03, 3.7439e-02,
-1.3838e-02, -3.6717e-02, -1.3899e-02, -1.4498e-02, 1.9595e-02,
-1.3527e-02, -2.5837e-02, 2.6661e-02, 4.4299e-02, 4.7894e-02,
-4.3256e-02, 3.0318e-02, 4.5738e-02, -2.5899e-02, 1.0935e-02,
-3.1078e-02, 3.5909e-02, -3.0156e-02, -3.8629e-02, -2.2232e-02,
2.9251e-02, 2.9346e-02, -2.2638e-02, -4.5489e-02, -2.4714e-02,
-6.5773e-05, 1.1286e-02, 2.4078e-02, 1.2962e-02, -1.4373e-02,
-3.6512e-02, -2.0977e-02, 3.0408e-02, -9.6226e-02, -2.7621e-02,
3.9301e-02, 4.6311e-02, 1.9279e-02, -1.8929e-02, 2.3344e-02,
5.2925e-02, -3.1078e-02, 4.0403e-02, -3.3261e-02, 3.8229e-02,
-3.3765e-02, -2.4055e-02, 3.1440e-02, 6.6322e-04, 3.4222e-02,
1.0625e-02, 8.2909e-03, 2.2984e-02, -1.6617e-02, 3.8297e-02,
-4.1747e-02, 3.4534e-02, -3.2529e-02, 4.3060e-02, 1.5835e-02,
1.1549e-02, -1.4029e-02, -7.6129e-02, -4.1971e-02, 4.5231e-02,
-3.3178e-02, -3.6076e-02, 9.6396e-02, 2.8793e-02, 4.7930e-02,
-1.2393e-02, -2.4239e-02, 4.2402e-02, -1.4571e-02, -1.8277e-02,
2.9203e-02, -1.7407e-02, -3.4116e-02, 2.6181e-02, 1.9214e-02,
1.6708e-02, 3.7283e-02, 3.2219e-02, -3.8511e-02, -3.6899e-02,
7.5839e-04, -1.2558e-02, -5.4088e-04, -1.5350e-02, -1.5896e-03,
8.2819e-03, 1.3798e-02, -4.2585e-02, -3.4217e-02, -1.7687e-02,
-1.3822e-02, 4.4292e-02, 1.0736e-02, -3.2278e-02, -1.2358e-02,
-3.3717e-02, 4.6526e-02, -2.0516e-02, 4.4928e-02, -4.8296e-02,
-3.4864e-02, -1.7273e-02, 1.5343e-02, -2.8694e-02, -1.7148e-03,
-4.8091e-02, -1.7803e-02, 6.6503e-03, -9.2686e-03, 4.5377e-02,
-1.3754e-02, -1.4563e-02, 2.7921e-02, -3.3957e-02, -4.8465e-02,
-2.4496e-02, 3.2424e-02, 1.0412e-02, -3.2561e-02, 2.9595e-02,
-2.9889e-02, 3.4239e-02, -2.8025e-02, -7.2978e-02, 4.5231e-02,
-3.2989e-02, 4.8832e-02, 9.6396e-02, -3.5873e-02, 4.5693e-02,
-1.5511e-02, 3.0311e-02, 7.2016e-02, -2.3770e-02, -1.8277e-02,
-1.3822e-02, -1.5371e-02, 1.7146e-03, -1.2413e-02, -1.6326e-02,
-2.7728e-02, -3.3387e-02, -3.5748e-02, -1.5689e-02, 2.4184e-02,
2.5512e-02, -1.4346e-02, 3.2656e-02, 7.4692e-04, 3.7456e-03,
9.2719e-03, 8.6953e-03, 1.6176e-02, 3.1407e-02, 3.3655e-02,
-4.8488e-02, -4.3406e-02, -3.1243e-02, -3.5053e-02, 1.1376e-02,
-3.2017e-02, -9.6133e-02, -3.9478e-02, 6.3121e-02, -4.4296e-02,
7.4575e-03, 3.7374e-02, 3.1938e-02, -3.0728e-02, -2.6533e-02,
-3.2724e-02, -1.0361e-02, 1.3341e-02, 4.1251e-02, -2.4932e-02,
1.6936e-02, -1.8192e-02, 1.7608e-02, -3.2066e-02, 2.3917e-02,
8.1959e-03], grad_fn=SelectBackward0)

In [36]: def get_default_device():
return torch.device('cpu')

def to_device(data, device):
if isinstance(data, (list,tuple)):
return [to_device(x, device) for x in data]
return data.to(device, non_blocking=True)

class DeviceDataLoader():
def __init__(self, dl, device):
self.dl = dl
self.device = device

def __iter__(self):
for b in self.dl:
yield to_device(b, self.device)

def __len__(self):
return len(self.dl)

In [37]: device = get_default_device()
device

Out[37]: device(type='cpu')

In [38]: train_dl = DeviceDataLoader(DataLoader(train_ds, batch_size=2), device)
val_dl = DeviceDataLoader(val_dl, device)
to_device(model, device);

In [40]: from torch.utils.tensorboard import SummaryWriter
import tensorflow as tf
tf.compat.v1.disable_eager_execution
@torch.no_grad()
def evaluate(model, val_loader):
model.eval()
outputs = [model.validation_step(batch) for batch in val_loader]
return model.validation_epoch_end(outputs)

def fit(epochs, lr, model, train_loader, val_loader, opt_func=torch.optim.SGD):
history = []
optimizer = opt_func(model.parameters(), lr)
for epoch in range(epochs):
# Training phase
model.train()
train_losses = []
for batch in train_loader:
loss = model.training_step(batch)
writer = SummaryWriter()
writer.add_scalar('Loss', loss, epoch)
train_losses.append(loss)
loss.backward()
optimizer.step() # validation phase
# validation phase
result = evaluate(model, val_loader)
result['train_loss'] = torch.stack(train_losses).mean().item()
model.eval()
model_epoch_end(epoch, result)
history.append(result)
return history

In [46]: model = to_device(StanfordCarsModel(), device)

In [47]: evaluate(model, val_dl)

Out[47]: ('val_loss': 5.289313248026855, 'val_acc': 0.0629296875)

In [48]: num_epochs = 10
opt_func = torch.optim.Adam
lr = 5e-2

In [49]: jvian.reset()
jvian.log_hyperparams({
'run_epochs': num_epochs,
'opt_func': opt_func.__name__,
'batch_size': batch_size,
'lr': lr,
})
[jvian] Hyperparams logged.

In [50]: history = fit(num_epochs, lr, model, train_dl, val_dl, opt_func)
Epoch [0], train_loss: 170070.6466, val_loss: 5.3159, val_acc: 0.0049
Epoch [1], train_loss: 5.2997, val_loss: 5.2099, val_acc: 0.0226
Epoch [2], train_loss: 5.2929, val_loss: 5.2083, val_acc: 0.0228
Epoch [3], train_loss: 5.2929, val_loss: 5.2961, val_acc: 0.0226
Epoch [4], train_loss: 5.5231, val_loss: 5.2059, val_acc: 0.0226
Epoch [5], train_loss: 5.2960, val_loss: 5.2966, val_acc: 0.0228
Epoch [6], train_loss: 5.2889, val_loss: 5.2952, val_acc: 0.0226
Epoch [7], train_loss: 5.2897, val_loss: 5.2966, val_acc: 0.0228
Epoch [8], train_loss: 5.2919, val_loss: 5.2829, val_acc: 0.0226
Epoch [9], train_loss: 5.2919, val_loss: 5.2829, val_acc: 0.0226

In [51]: jvian.reset()
jvian.log_hyperparams({
'run_epochs': num_epochs,
'opt_func': opt_func.__name__,
'batch_size': batch_size,
'lr': lr,
})
[jvian] Hyperparaas logged.

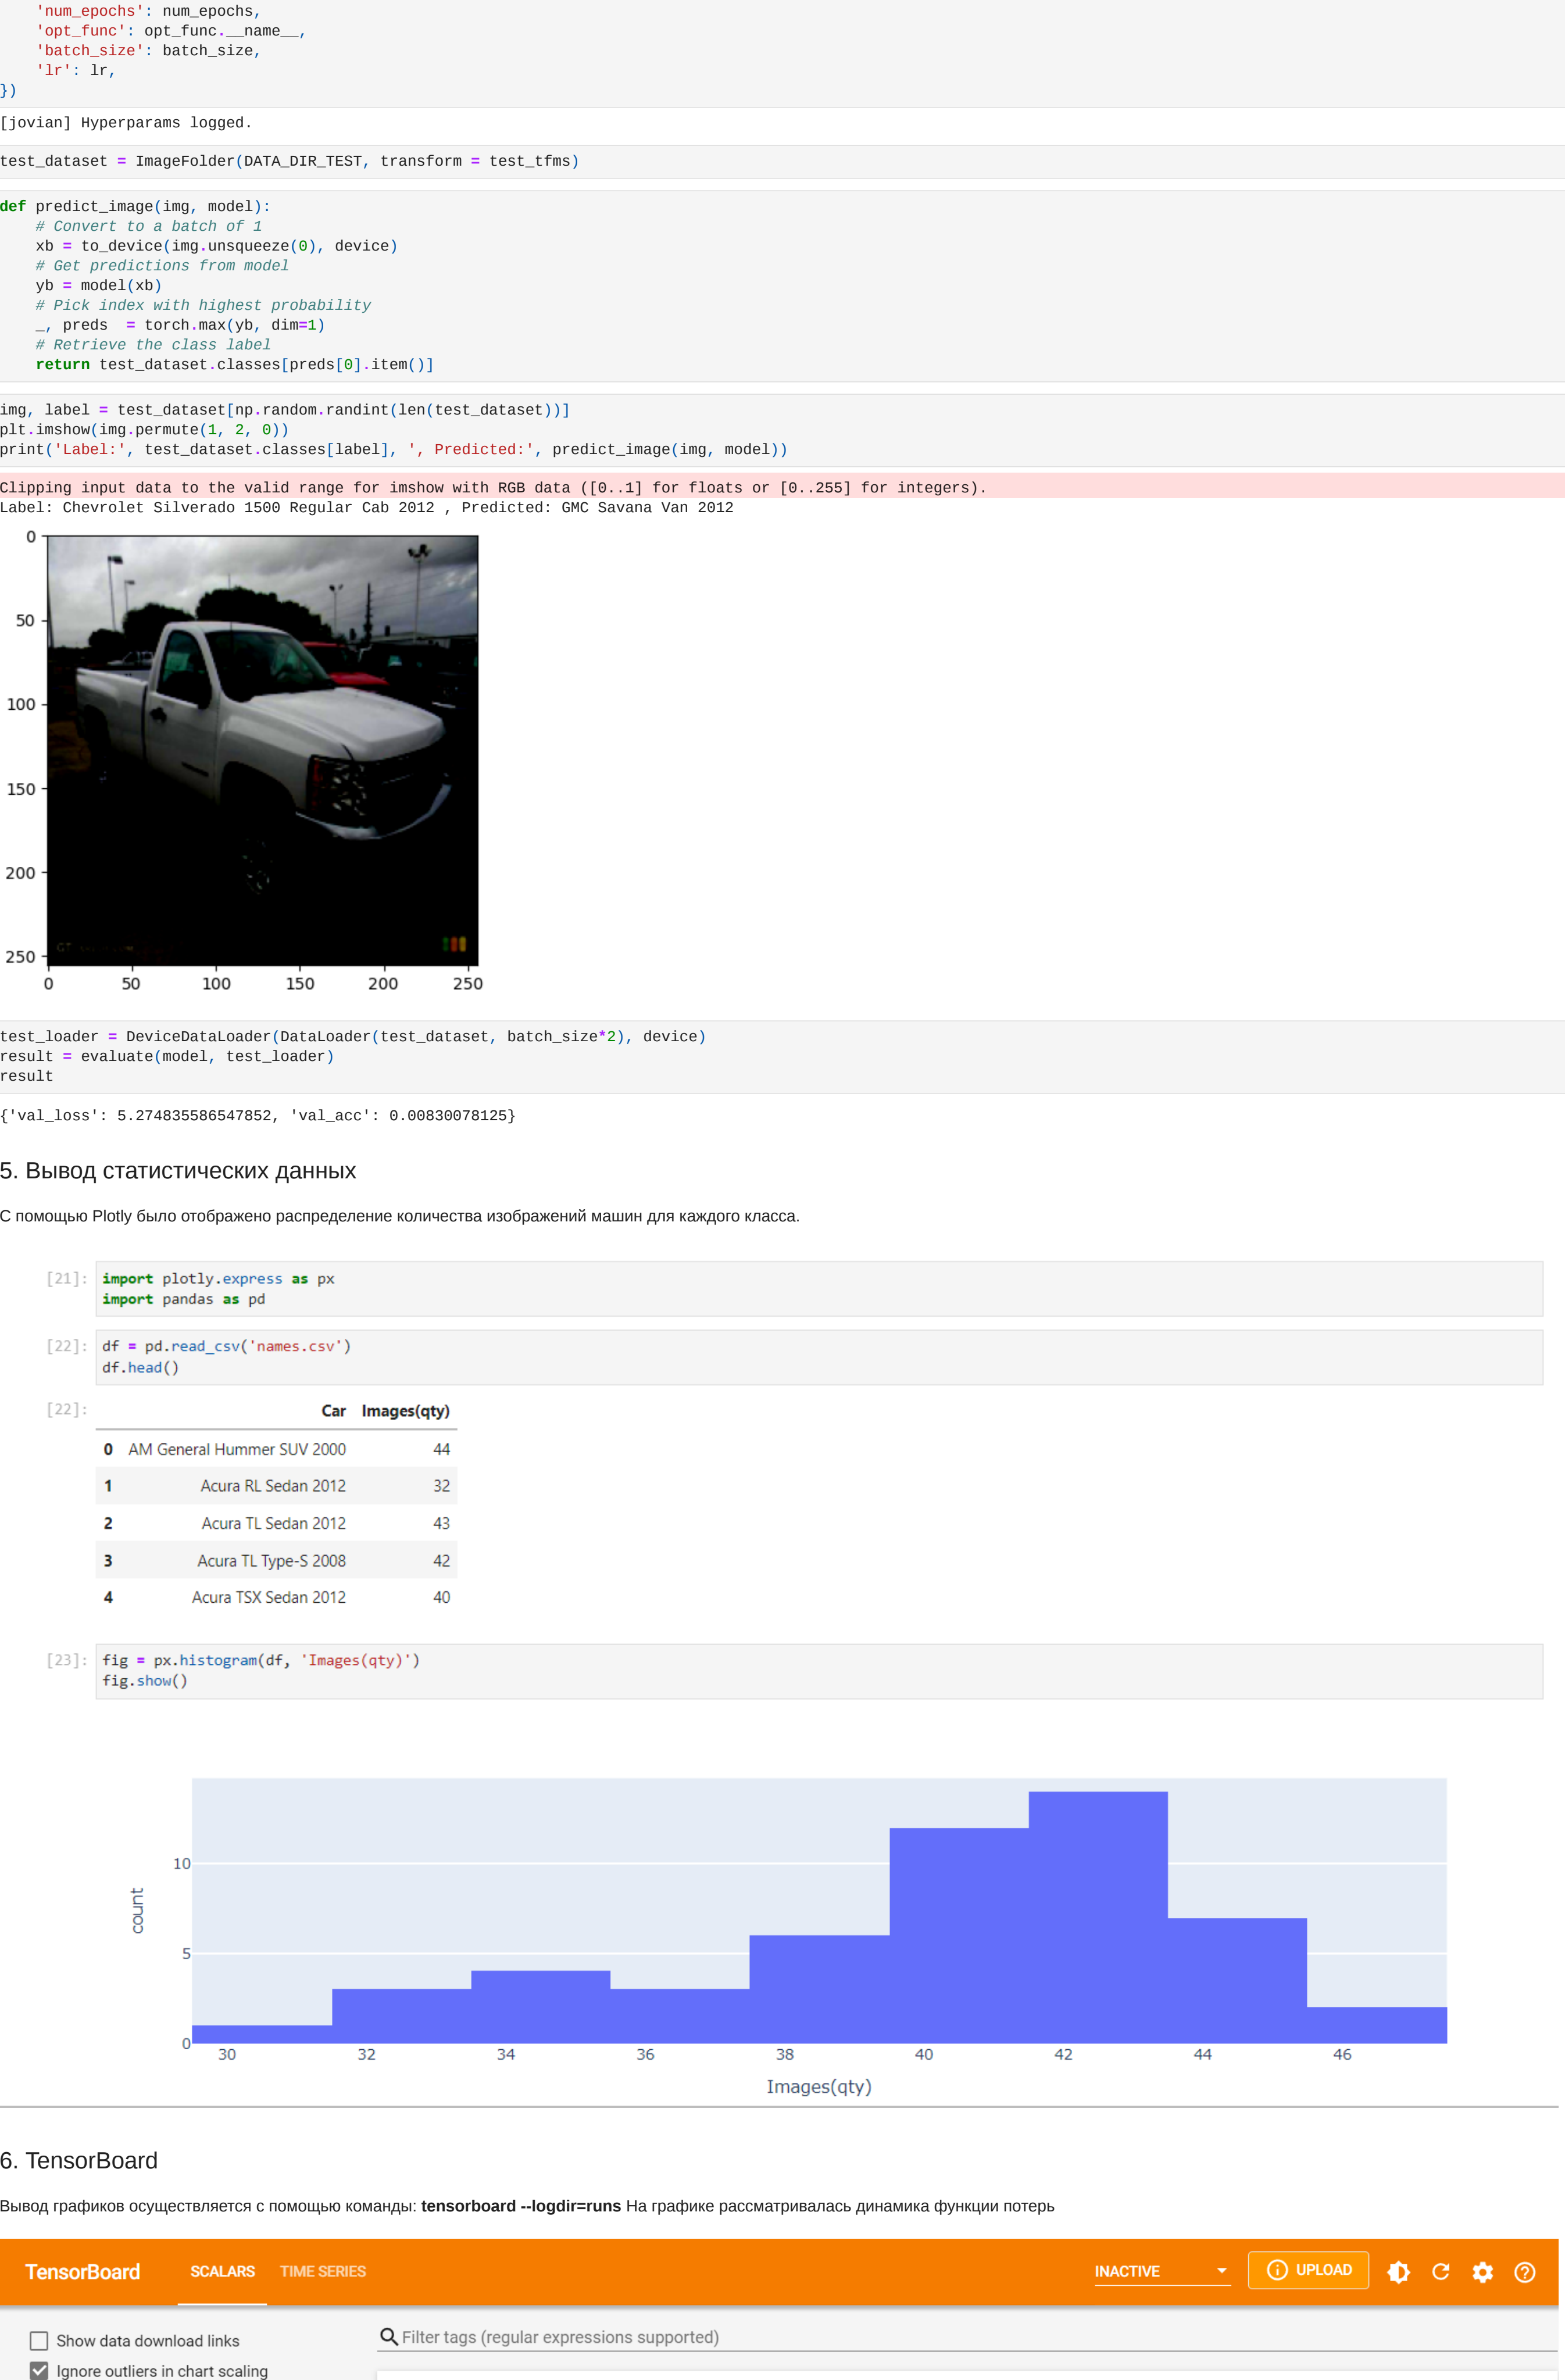
In [52]: test_dataset = ImageFolder(DATA_DIR_TEST, transform = test_tfms)

In [53]: def predict_image(img, model):
# Convert to a batch of 1
xb = to_device(img.unsqueeze(0), device)
# Get predictions from model
yb = model(xb)
# Pick index with highest probability
_, preds = torch.max(yb, dim=1)
# Return the class label
return test_dataset.classes[preds[0].item()]

In [56]: img, label = test_dataset[np.random.randint(len(test_dataset))]
plt.imshow(img.permute(1, 2, 0))
print('Label:', test_dataset.classes[label], ', Predicted:', predict_image(img, model))

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Label: Chevrolet Silverado 1500 Regular Cab 2012, Predicted: GMC Savana Van 2012
```



6. TensorBoard

Вывод графиков осуществляется с помощью команды `tensorboard --logdir=runs` на графике рассматривалась динамика функции потерь

